

Class 10

Yuntian Zhu (PID: A17816597)

Table of contents

Data Import	1
Quick overview of the dataset	3
Winpercent and Pricepercent	11
Exploring the correlation structure	13
Principal Component Analysis	14

Today we will do a mini project about Halloween

Our data come from the 538 website and is available as a CSV file

Data Import

```
candy <- read.csv("candy-data.csv", row.names = "competitorname")  
head(candy)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	crispedrice	wafers
100 Grand	1	0	1	0	0	0	0	1
3 Musketeers	1	0	0	0	0	1	0	0
One dime	0	0	0	0	0	0	0	0
One quarter	0	0	0	0	0	0	0	0
Air Heads	0	1	0	0	0	0	0	0
Almond Joy	1	0	0	0	1	0	0	0
	hard	bar	pluribus	sugarpercent	pricepercent	winpercent		
100 Grand	0	1	0	0.732	0.860	66.97173		
3 Musketeers	0	1	0	0.604	0.511	67.60294		
One dime	0	0	0	0.011	0.116	32.26109		
One quarter	0	0	0	0.011	0.511	46.11650		

Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

```
flextable::flextable(head(candy, 10))
```

chocolate	fruity	caramel	peanut	almond	dyno	nougat	crisped	rice	wafer	hard	bar	pluribus	s
1	0	1	0	0	0	1	1	0	0	0	1	0	0
1	0	0	0	0	1	0	0	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	1	0	0	0	0	0	0	0	0	0	0	0	0
1	0	0	1	0	0	0	0	0	0	0	1	0	0
1	0	1	1	1	1	0	0	0	0	0	1	0	0
0	0	0	1	0	0	0	0	0	0	0	0	0	1
0	0	0	0	0	0	0	0	0	0	0	0	0	1
0	1	1	0	0	0	0	0	0	0	0	0	0	0

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

[1] 85

There are 85 types of candies in this dataset.

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

[1] 38

There are 38 fruity candy types in the dataset.

Q3 What is your favorite candy in the dataset and what is it's winpercent value?

My favorite candy in the dataset is Almond Joy

```
candy["Almond Joy", "winpercent"]
```

```
[1] 50.34755
```

Almond Joy's winpercent is 50.35%.

Q4. What is the winpercent value for "Kit Kat"?

```
candy["Kit Kat", "winpercent"]
```

```
[1] 76.7686
```

Kitt Kat's winpercent is 76.77%

Q5. What is the winpercent value for "Tootsie Roll Snack Bars"?

```
candy["Tootsie Roll Snack Bars", "winpercent"]
```

```
[1] 49.6535
```

Tootsie Roll Snack Bars's winpercent is 49.65%

Quick overview of the dataset

```
library (skimr)  
skim(candy)
```

Table 2: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12

Group variables	None
-----------------	------

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmond	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

The winpercent column has a very different scale, compared to the majority of the other columns.

Q7. What do you think a zero and one represent for the candy\$chocolate column?


The zero means that candy does not contain chocolate. The one mean that the candy contain chocolate.

Q8. Plot a histogram of winpercent values

```
library(ggplot2)

# Plot histogram
ggplot(candy, aes(x = winpercent)) +
  geom_histogram(bins = 20, fill = "steelblue", color = "white") +
  labs(
    title = "Distribution of Candy Win Percentages",
    x = "Win Percent",
    y = "Count"
  )
```

Distribution of Candy Win Percentages




Q9. Is the distribution of winpercent values symmetrical?

```
library(ggplot2)

# Plot histogram
ggplot(candy, aes(x = winpercent)) +
  geom_density(bins = 20, fill = "steelblue", color = "white") +
  labs(
    title = "Distribution of Candy Win Percentages",
    x = "Win Percent",
    y = "Count"
  )
```

Warning in geom_density(bins = 20, fill = "steelblue", color = "white"):
Ignoring unknown parameters: `bins`

Distribution of Candy Win Percentages



The distribution of `winpercent` is not symmetrical. Instead, it is skewed to the right in terms of statistical terminology.

Q10. Is the center of the distribution above or below 50%?

```
mean(candy$winpercent)
```

```
[1] 50.31676
```

```
summary(candy$winpercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
22.45	39.14	47.83	50.32	59.86	84.18

If we use the mean as the measure of the center, it is indeed above 50%.

However, since the data is asymmetrical, median is a better measure of the center, and it is below 50%.

Q11. On average is chocolate candy higher or lower ranked than fruity candy?

```

# 1. Find all chocolate candy
# 2. Find their winpercent values
# 3. Calculate the mean of these values
# 4-6 Do the same for fruity candy
# 7. Compare mean winpercents of chocolate vs fruity
# 8. Pick the highest as the winner

choc inds <- candy$chocolate == 1
choc.win <- candy[choc inds, ]$winpercent
choc.mean <- mean(choc.win)
choc.mean

```

[1] 60.92153

```

fruity inds <- candy$fruity == 1
fruity.win <- candy[fruity inds, ]$winpercent
fruity.mean <- mean(fruity.win)
fruity.mean

```

[1] 44.11974

On average, chocolate candies are ranked higher than fruity candy.

Q12. Is this difference statistically significant?

```
t.test(choc.win, fruity.win)
```

Welch Two Sample t-test

```

data: choc.win and fruity.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974

```

The t-test results in a p value that is extremely small (2.871e-8), which is significantly below the commonly used cutoff (5%). Therefore, the difference is statistically significant.

Q13. What are the five least liked candy types in this set?

```
library(tidyverse)

-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr     1.1.4     v readr     2.1.5
v forcats   1.0.1     v stringr   1.5.2
v lubridate 1.9.4     v tibble    3.3.0
v purrr     1.1.0     v tidyr    1.3.1
-- Conflicts -----
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become non-conflicting
```

```
library(tibble)

candy %>%
  rownames_to_column("competitorname") %>%      # restore candy names as a column
  arrange(winpercent) %>%                         # sort by winpercent ascending
  select(competitorname, winpercent) %>%          # select only these columns
  head(5)                                         # show 5 least liked candies
```

	competitorname	winpercent
1	Nik L Nip	22.44534
2	Boston Baked Beans	23.41782
3	Chiclets	24.52499
4	Super Bubble	27.30386
5	Jawbusters	28.12744

The five least liked candy types are Nik L Nip, Boston Baked Beans, Chiclets, Super Bubble and Jawbusters.

Q14. What are the top 5 all time favorite candy types out of this set?


```
candy %>%
  rownames_to_column("competitorname") %>%      # restore candy names as a column
  arrange(winpercent) %>%                         # sort by winpercent ascending
  select(competitorname, winpercent) %>%          # select only these columns
  tail(5)                                         # show 5 most liked candies
```

	competitorname	winpercent
81	Snickers	76.67378
82	Kit Kat	76.76860
83	Twix	81.64291
84	Reese's Miniatures	81.86626
85	Reese's Peanut Butter cup	84.18029

The five most liked candy types are Snickers, Kit Kat, Twix, Reese's Miniatures and Reese's Peanut Butter Cup.


Q15. Make a first barplot of candy ranking based on winpercent values

```
ggplot(candy) +  
  aes(winpercent, rownames(candy)) +  
  geom_col()
```



Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?


```
ggplot(candy) +  
  aes(winpercent, reorder(rownames(candy), winpercent)) +  
  geom_col()
```



Add some color based on the “type of candy”

```
my_cols <- rep("black", nrow(candy))
my_cols[as.logical((candy$chocolate))] <- "chocolate"
my_cols[as.logical((candy$fruity))] <- "brown"
my_cols[as.logical((candy$bar))] <- "pink"
```

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col(fill = my_cols)
```



Q17. What is the worst ranked chocolate candy?

The worst ranked chocolate candy is Sixlets.

Q18. What is the best ranked fruity candy?


The best ranked fruity candy is Starburst.

Winpercent and Pricepercent

A plot with both variables/columns winpercent and pricepercent

```
library(ggrepel)
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

We can answer this question by calculating the ratio between winpercent and pricepercent.

```
candy %>%
  mutate(
    bang_for_buck = winpercent / pricepercent,
    name = rownames(.)
  ) %>%
  arrange(desc(bang_for_buck)) %>%
  slice(1)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Tootsie Roll Midgies	1	0	0	0	0	0
	crispedrice	wafer	hard bar	pluribus	sugar	percent
Tootsie Roll Midgies	0	0	0	1	0.174	
	pricepercent	winpercent	bang_for_buck			name
Tootsie Roll Midgies	0.011	45.73675	4157.886	Tootsie Roll Midgies		

Therefore, the candy type is Tootsie Roll Midgies.

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
candy %>%
  arrange(desc(pricepercent)) %>%
  head(5) %>%
  mutate(name = rownames(.)) %>%
  arrange(winpercent)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat	
Nik L Nip	0	1	0	0	0	0	
Ring pop	0	1	0	0	0	0	
Nestle Smarties	1	0	0	0	0	0	
Hershey's Milk Chocolate	1	0	0	0	0	0	
Hershey's Krackel	1	0	0	0	0	0	
	crisped	rice	wafer	hard	bar	pluribus	sugarpercent
Nik L Nip	0	0	0	1			0.197
Ring pop	0	1	0	0			0.732
Nestle Smarties	0	0	0	1			0.267
Hershey's Milk Chocolate	0	0	1	0			0.430
Hershey's Krackel	1	0	1	0			0.430
	pricepercent	winpercent					name
Nik L Nip	0.976	22.44534					Nik L Nip
Ring pop	0.965	35.29076					Ring pop
Nestle Smarties	0.976	37.88719					Nestle Smarties
Hershey's Milk Chocolate	0.918	56.49050	Hershey's Milk Chocolate				
Hershey's Krackel	0.918	62.28448					Hershey's Krackel

The 5 most expensive candy types are Nik L Nip, Ring pop, Nestle Smarties, Hershey's Milk Chocolate and Hershey's Krackel. The least popular one is Nik L Nip.


Exploring the correlation structure

Now that we've explored the dataset a little, we'll see how the variables interact with one another. We'll use correlation and view the results with the corrplot package to plot a correlation matrix.

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
cij <- cor(candy)
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Fruity and chocolate are the two most anti-correlated variables.

Q23. Similarly, what two variables are most positively correlated?

Each candy type is positively correlated with itself (coefficient = 1). Chocolate is highly correlated with bar and winpercent.

Principal Component Analysis

The function to use is called `prcomp()`. It has an optional `scale` argument.

```
pca <- prcomp(candy, scale = TRUE)
summary(pca)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

Our main PCA result figure


```
ggplot(pca$x) +
  aes(PC1, PC2, label = rownames(pca$x)) +
  geom_point(col = my_cols) +
  geom_text_repel(size=3.3, col=my_cols, max.overlaps = 7)
```

Warning: ggrepel: 51 unlabeled data points (too many overlaps). Consider increasing max.overlaps



We should also examine the variable “loadings” or contributions of the original variables to the new PCs

```
ggplot(pca$rotation) +
  aes(PC1, rownames(pca$rotation)) +
  geom_col()
```



Interactive plots that can be zoomed on and “brushed” over can be made with the **plotly** package. But this does not work in PDF.

Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

Fruity, hard and pluribus. This makes sense to me because all the brown points (the color of fruity) are on the right part of the pca plot.