

# Class12

Yuntian Zhu (PID: A17816597)

## Table of contents

Background . . . . .	1
Data import . . . . .	1
Toy differential gene expression . . . . .	3
DESeq2 analysis . . . . .	9
Volcano Plot . . . . .	11
Save our results . . . . .	12

## Background

Today we will analyze some RNA-seq data on the effects of a common steroid on airway smooth muscle cells.

## Data import

```
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
head(counts)
```

	SRR1039508	SRR1039509	SRR1039512	SRR1039513	SRR1039516
ENSG00000000003	723	486	904	445	1170
ENSG00000000005	0	0	0	0	0
ENSG00000000419	467	523	616	371	582
ENSG00000000457	347	258	364	237	318
ENSG00000000460	96	81	73	66	118
ENSG00000000938	0	0	1	0	2
	SRR1039517	SRR1039520	SRR1039521		

```
ENSG000000000003      1097      806      604
ENSG000000000005        0         0         0
ENSG000000000419      781      417      509
ENSG000000000457      447      330      324
ENSG000000000460       94      102       74
ENSG000000000938       0         0         0
```

```
head(metadata)
```

```
      id    dex celltype geo_id
1 SRR1039508 control   N61311 GSM1275862
2 SRR1039509 treated   N61311 GSM1275863
3 SRR1039512 control   N052611 GSM1275866
4 SRR1039513 treated   N052611 GSM1275867
5 SRR1039516 control   N080611 GSM1275870
6 SRR1039517 treated   N080611 GSM1275871
```

Q1. How many genes are in this dataset?

```
nrow(counts)
```

```
[1] 38694
```

There are 38694 genes in this dataset.

Q. How many different experiments (columns in counts or rows in metadata) are there?

```
ncol(counts)
```

```
[1] 8
```

Q2. How many ‘control’ cell lines do we have?

```
sum(metadata$dex == "control")
```

```
[1] 4
```

We have 4 control cell lines.

## Toy differential gene expression

To start our analysis, let's calculate the mean counts of the counts for all genes in the "control" experiments

1. Extract all "control" columns from the `counts` object
2. Calculate the mean for all rows (i. e. genes) of these "control" columns
- 3 and 4. Do the same thing for "treated" groups
5. Compare these `control.mean` and `treated.mean` values.

Method 1:

```
control <- metadata[metadata[, "dex"] == "control",]  
control.counts <- counts[, control$id]  
control.mean <- rowSums(control.counts) / 4  
head(control.mean)
```

ENSG00000000003	ENSG00000000005	ENSG000000000419	ENSG000000000457	ENSG000000000460
900.75	0.00	520.50	339.75	97.25
ENSG000000000938				
0.75				

Q3. How would you make the above code in either approach more robust? Is there a function that could help here?

The method above finds the mean by first calculate the summation and divide it by 4 (the number of samples), which is not very robust. The primary reason is that there might be a different sample number and we need to change 4 to another number by ourselves. We can improve the code by using the function `rowMeans()`. By using `rowMeans()`, the mean is calculated automatically across however many samples are present, without requiring us to hard-code the sample count. This not only makes the code more flexible and maintainable, but also reduces the risk of errors when working with different datasets or experimental designs.

```
control inds <- metadata$dex == "control"  
control.count <- counts[, control.inds]  
control.mean <- rowMeans(control.count)  
head(control.mean)
```

ENSG00000000003	ENSG00000000005	ENSG000000000419	ENSG000000000457	ENSG000000000460
900.75	0.00	520.50	339.75	97.25
ENSG000000000938				
0.75				

Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug treated samples and assign to a labeled vector called treated.mean)

```
treated inds <- metadata$dex == "treated"  
treated.count <- counts[, treated inds]  
treated.mean <- rowMeans(treated.count)  
head(treated.mean)
```

```
ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460  
658.00 0.00 546.00 316.50 78.75  
ENSG000000000938  
0.00
```

Store these all together for the ease of bookkeeping as meancounts


```
meancounts <- data.frame(control.mean, treated.mean)  
head(meancounts)
```

	control.mean	treated.mean
ENSG000000000003	900.75	658.00
ENSG000000000005	0.00	0.00
ENSG000000000419	520.50	546.00
ENSG000000000457	339.75	316.50
ENSG000000000460	97.25	78.75
ENSG000000000938	0.75	0.00

Q5 (a). Create a scatter plot showing the mean of the treated samples against the mean of the control samples.

Let us make a plot of control vs. treated mean values


```
plot(meancounts)
```



Q5 (b). You could also use the ggplot2 package to make this figure producing the plot below. What geom\_?() function would you use for this plot?

Let us make another plot using ggplot2. We should use `geom_point()` here.

```
library(ggplot2)
ggplot(meancounts) +
  aes(control.mean, treated.mean, alpha = 0.2) +
  geom_point()
```




Q6. Try plotting both axes on a log scale. What is the argument to plot() that allows you to do this?

Make this a log log plot, so that the plot gets less crowded. The argument that we should use here is `log = "xy"`.

```
plot(meancounts, log = "xy")
```

Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted from logarithmic plot

Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted from logarithmic plot



We often talk about metrics like log2FC. Let's calculate the log2 fold change for our treated over control mean counts.

```
meancounts$log2fc <-
log2(meancounts$treated.mean/
  meancounts$control.mean)
head(meancounts)
```

	control.mean	treated.mean	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG000000000005	0.00	0.00	NaN
ENSG000000000419	520.50	546.00	0.06900279
ENSG000000000457	339.75	316.50	-0.10226805
ENSG000000000460	97.25	78.75	-0.30441833
ENSG000000000938	0.75	0.00	-Inf

A common “rule of thumb” is a log2 fold change cutoff of +2 and -2 to call genes “upregulated” or “downregulated”.

First, we can find the number of upregulated genes

```
sum(meancounts$log2fc >= +2, na.rm = T)
```

```
[1] 1910
```

Then, we can find the number of downregulated genes

```
sum(meancounts$log2fc >= -2, na.rm = TRUE)
```

```
[1] 23046
```

We can make the data more polished by removing the genes that are not even expressed, so that we will not have things like NaN or -Inf.

```
zero.vals <- which(meancounts[,1:2]==0, arr.ind=TRUE)

to.rm <- unique(zero.vals[,1])
mycounts <- meancounts[-to.rm,]
head(mycounts)
```

	control.mean	treated.mean	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG000000000419	520.50	546.00	0.06900279
ENSG000000000457	339.75	316.50	-0.10226805
ENSG000000000460	97.25	78.75	-0.30441833
ENSG000000000971	5219.00	6687.50	0.35769358
ENSG000000001036	2327.00	1785.75	-0.38194109

Q7. What is the purpose of the `arr.ind` argument in the `which()` function call above? Why would we then take the first column of the output and need to call the `unique()` function?

The purpose of `arr.ind` is to get row and column coordinates instead of just linear positions, because normally, `which()` returns just the indices of TRUE values in a vector. The purpose of taking the first column is to find which rows (genes) have a zero in either `control.mean` or `treated.mean`. The first column of `zero.vals` corresponds to the row index of each zero entry, so that we can remove any genes that have zero counts in any samples. The purpose of calling the `unique()` function is that A single row (gene) might appear multiple times if both `control.mean` and `treated.mean` are zero, or if repeated zeros are found. To avoid removing the same row multiple times, we take only unique row indices.

Q8. Using the up.ind vector above can you determine how many up regulated genes we have at the greater than 2 fc level?

```
up.ind <- mycounts$log2fc > 2  
sum(up.ind)
```

```
[1] 250
```

We have 250 upregulated genes at the greater than 2 fc level.

Q9. Using the down.ind vector above can you determine how many down regulated genes we have at the greater than 2 fc level?

```
down.ind <- mycounts$log2fc < (-2)  
sum(down.ind)
```

```
[1] 367
```

We have 367 downregulated genes at the greater than 2 fc level.

Q10. Do you trust these results? Why or why not?

Not really. The results only reflect whether the log2FC of a certain gene is greater than the 2 cutoff or less than the -2 cutoff. In order to make the result more trustable, we should focus on the genes that show statistically significant change. In other words, we should also consider the p-values in addition to log2FC. Then, we can focus on the genes that not only have log2FC  $> 2$  or log2FC  $< -2$  but also have a small p-value showing the statistical significance.

## DESeq2 analysis

Let's do this analysis properly and find genes that have a statistically significant change.

```
library(DESeq2)
```

For DESeq analysis, we need three things:

- count values (`countData`)
- metadata telling us about the columns in `countData` (`colData`)
- design of the experiment (i. e. what do you want to compare)

Our first function from DESeq2 will set up the input for our analysis by storing all these three things together.

```
dds <- DESeqDataSetFromMatrix(countData = counts,
                                colData = metadata,
                                design = ~dex)
```

converting counts to integer mode

Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in  
design formula are characters, converting to factors

The main function in DESeq2 that runs the differential expression analysis is called `DESeq()`

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

```
res <- results(dds)
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 6 columns

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
ENSG000000000003	747.194195	-0.350703	0.168242	-2.084514	0.0371134
ENSG000000000005	0.000000	NA	NA	NA	NA

```

ENSG00000000419 520.134160      0.206107  0.101042  2.039828 0.0413675
ENSG00000000457 322.664844      0.024527  0.145134  0.168996 0.8658000
ENSG00000000460 87.682625      -0.147143  0.256995 -0.572550 0.5669497
ENSG00000000938 0.319167      -1.732289  3.493601 -0.495846 0.6200029
          padj
<numeric>
ENSG00000000003 0.163017
ENSG00000000005   NA
ENSG00000000419 0.175937
ENSG00000000457 0.961682
ENSG00000000460 0.815805
ENSG00000000938   NA

```


## Volcano Plot

This is a common summery result figure from these types of experiments an plot the log2 fold-change vs the adjusted p-value

```

plot(res$log2FoldChange, -log(res$padj))
abline(v = c(-2, 2), col = "red")
abline(h = -log(0.05), col = "red")

```



## Save our results

```
write.csv(res, file = "myresults.csv")
```